

# Q1 Model


## ◆ Model

- Describe the model architecture and how it works on text summarization.

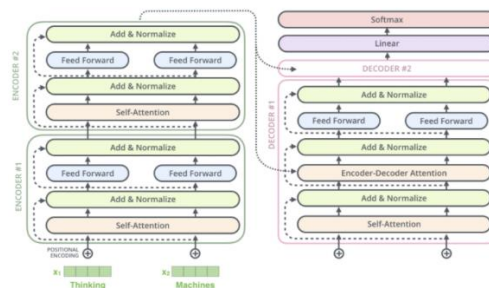
BERT 主要做 NLU 任務，對於 NLG 有點力不從心，但是 T5 用一種大一統的思想把 NLU 和 NLG 任務一起解決了，這一種大一統的思想就是：把所有的 NLP 任務都轉化成了文字到文字 Text2Text 格式的任務。而 MT5 就是 Multilingual 基於 T5 的架構。

T5 是基於 transformer - Encoder-Decoder 架構如下圖

```
{
  "architectures": [
    "MT5ForConditionalGeneration"
  ],
  "d_ff": 1024,
  "d_kv": 64,
  "d_model": 512,
  "decoder_start_token_id": 0,
  "dropout_rate": 0.1,
  "eos_token_id": 1,
  "feed_forward_proj": "gated-gelu",
  "initializer_factor": 1.0,
  "is_encoder_decoder": true,
  "layer_norm_epsilon": 1e-06,
  "model_type": "mt5",
  "num_decoder_layers": 8,
  "num_heads": 6,
  "num_layers": 8,
  "pad_token_id": 0,
  "relative_attention_num_buckets": 32,
  "tie_word_embeddings": false,
  "tokenizer_class": "T5Tokenizer",
  "vocab_size": 250112
}
```



其整體架構如下：



MT5 可以做許多不同種類的任務，像翻譯任務，分類任務，回歸任務，而 summarization 任務就是 seq2seq 任務。

## ◆ Preprocessing

- Describe your preprocessing (e.g. tokenization, data cleaning and etc.)



將 input content size (maintext) 設為 1024，當 maintext 大於 1024 時，刪掉超過長度的字，不到 1024 的則做 padding，在將字轉為 vocab 裡的 id。

mT5 涵蓋了 101 種語言，總詞表(vocab)有 25 萬，而且它採用的 T5 結構的 Softmax 還不共享引數，這就導致了 Embedding 層佔用了相當多的引數量，比如 mT5 small 的參數量為 3 億，其中 Embedding 相關的就佔了 2.5 億，而大部分的參數我們都用不到，因此要精簡 Embedding，其原理就是需要在兩個 Embedding 矩陣中刪除不需要的行就行了，關鍵在於如何決定要保留的 token，以及如何得到一個精簡後的 sentencepiece 模型。決定要保留的 token，簡單來想就是把中文的 token 保留下來，但是也不只是中文，英文的也要保留一部分。


## Q2 Training

### ◆ Hyperparameter

- Describe your hyperparameter you use and how you decide it.

當選擇小的 beam size 可以使產生的文字較符合主題，但是文法會比較不通順。

當選擇大的 beam size 可以使產生的句子文法較通順且回答較正確，但是較通用且無相關性。



A cartoon illustration of a person with dark hair and a blue shirt. Above their head is a speech bubble containing the text: "I mostly eat a fresh and raw diet, so I save on groceries".


Beam Size	Model Response
1	<i>I love to eat healthy and eat healthy</i>
2	<i>That is a good thing to have</i>
3	<i>I am a nurse so I do not eat raw food</i>
4	<i>I am a nurse so I am a nurse</i>
5	<i>Do you have any hobbies?</i>
6	<i>What do you do for a living?</i>
7	<i>What do you do for a living?</i>
8	<i>What do you do for a living?</i>

經過多次的測試，我最後使用 beam size = 7，並搭配 No\_repeat\_ngram\_size=3 來防止一直產生相同的詞。

max\_source\_length=2048  
max\_target\_length=128  
learning\_rate=3e-5  
train\_batch\_size=1  
gradient\_accumulation\_steps=16  
eval\_batch\_size=6

## ◆ Learning Curves

- Plot the learning curves (ROUGE versus training steps)




## Q3 Generation Strategies

◆ strategies

- Describe the detail of the following generation strategies:


- Greedy

選擇最大機率的字(argmax)



- Beam Search

選擇 k 個最大機率的 sequence 並且找出最佳(最大機率)的一句 sequence。



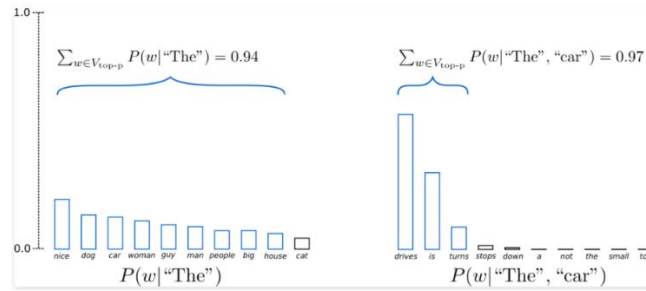
- Top-k Sampling

將挑選出  $K$  個最有可能的下一個單詞，並且僅在這  $K$  個下一個單詞之間重新為它們分配概率。

- **Top-p Sampling**

在 Top-p 採樣中，不是僅從最可能的 K 個單詞中採樣，而是從其累積概率超過一個閾值 的最小可能單詞集中進行選擇，然後將這組單詞重新分配概率。這樣，單詞集合的大小（也就是集合中單詞的數量）可以根據下一個單詞的概率分布動態地增加或減少。

如下所示:



- **Temperature** (不是一種 strategy，而是一種附加的參數可以搭配不同的 Generation Strategies 使用)

**Ex :**

Softmax temperature: applying a **temperature hyperparameter**  $\tau$  to the softmax

$$P(w_t) = \frac{e^{s_w/\tau}}{\sum_{w' \in V} e^{s_{w'}/\tau}}$$

Higher temperature:  $P(w_t)$  becomes more uniform  $\rightarrow$  more diversity  
 Lower temperature:  $P(w_t)$  becomes more spiky  $\rightarrow$  less diversity

## ◆ Hyperparameters

- Try at least 2 settings of each strategies and compare the result.
  - Top\_k (200) + top\_p (0.9) + temperature(0.9)

Epoch=20

Input size : 2048

Output size : 128

```
gen_kwargs = {
    "max_length": args.val_max_target_length if args is not None else config.max_length,
    # "num_beams": args.num_beams,
}

for step, batch in enumerate(tqdm(eval_dataloader)):
    with torch.no_grad():
        generated_tokens = accelerator.unwrap_model(model).generate(
            batch["input_ids"],
            attention_mask=batch["attention_mask"],
            **gen_kwargs,
            top_k=200,
            top_p=0.9,
            do_sample=True,
            no_repeat_ngram_size=2,
            length_penalty=1.0,
            # num_return_sequences=5,
            temperature=0.9
        )
```

Result:

```
{
  "rouge-1": {
    "f": 0.233834156220925,
    "p": 0.23770516832925132,
    "r": 0.24138568730958734
  },
  "rouge-2": {
    "f": 0.08343446488783317,
    "p": 0.08486256536017532,
    "r": 0.08652563689559623
  },
  "rouge-l": {
    "f": 0.20825253250619444,
    "p": 0.2131501573446441,
    "r": 0.2131291228248939
  }
}
```

- Top\_k (20) + top\_p (0.98) + temperature(0.95)

Epoch=20

Input size : 2048

Output size : 128

```
gen_kwargs = {
    "max_length": args.val_max_target_length if args is not None else config.max_length,
    # "num_beams": args.num_beams,
}

for step, batch in enumerate(tqdm(eval_dataloader)):
    with torch.no_grad():
        generated_tokens = accelerator.unwrap_model(model).generate(
            batch["input_ids"],
            attention_mask=batch["attention_mask"],
            **gen_kwargs,
            top_k=20,
            top_p=0.98,
            do_sample=True,
            no_repeat_ngram_size=2,
            length_penalty=1.0,
            # num_return_sequences=5,
            temperature=0.95
        )
```

Result:

```
{
  "rouge-1": {
    "f": 0.23535839680203713,
    "p": 0.2394467326381215,
    "r": 0.24292915867524267
  },
  "rouge-2": {
    "f": 0.08377209643883991,
    "p": 0.08490877035974002,
    "r": 0.08699999718388637
  },
  "rouge-l": {
    "f": 0.20798106454380164,
    "p": 0.21321483321684892,
    "r": 0.2125747874588796
  }
}
```

- Top\_k (200) + top\_p (0.9) + temperature(0.3)

Epoch=20

Input size : 2048

Output size : 128

```
gen_kwargs = {
    "max_length": args.val_max_target_length if args is not None else config.max_length,
    # "num_beams": args.num_beams,
}

for step, batch in enumerate(tqdm(eval_data_loader)):
    with torch.no_grad():
        generated_tokens = accelerator.unwrap_model(model).generate(
            batch["input_ids"],
            attention_mask=batch["attention_mask"],
            **gen_kwargs,
            top_k=200,
            top_p=0.9,
            do_sample=True,
            no_repeat_ngram_size=2,
            length_penalty=1.0,
            # num_return_sequences=5,
            temperature=0.3
        )
```

Result:

```
{
  "rouge-1": {
    "f": 0.2595305394176496,
    "p": 0.2692356949705075,
    "r": 0.2620393624302309
  },
  "rouge-2": {
    "f": 0.09864152097394568,
    "p": 0.1025161004487499,
    "r": 0.0998890773058013
  },
  "rouge-l": {
    "f": 0.2304447916811395,
    "p": 0.241074326424151,
    "r": 0.23025836914758732
  }
}
```



- Beams\_number (7) + ngram(3)

Epoch=20

Input size = 2048

Output size = 128

```
gen_kwargs = {
    "max_length": args.val_max_target_length if args is not None else config.max_length,
    "num_beams": args.num_beams,
}

for step, batch in enumerate(tqdm(eval_data_loader)):
    with torch.no_grad():
        generated_tokens = accelerator.unwrap_model(model).generate(
            batch["input_ids"],
            attention_mask=batch["attention_mask"],
            **gen_kwargs,
            no_repeat_ngram_size=3,
            # length_penalty=1.0
            num_return_sequences=7
        )
```

Result:

```
{
  "rouge-1": {
    "f": 0.2783196093622653,
    "p": 0.2854129922686215,
    "r": 0.2843418305349766
  },
  "rouge-2": {
    "f": 0.1164916516785287,
    "p": 0.11983329519416126,
    "r": 0.11913228357826665
  },
  "rouge-l": {
    "f": 0.24909060026246263,
    "p": 0.2572848723185751,
    "r": 0.2522693397227175
  }
}
```

- Beams\_number (5) + ngram(3)

Epoch=20

Input size = 2048

Output size = 128

```
gen_kwargs = {
    "max_length": args.val_max_target_length if args is not None else config.max_length,
    "num_beams": args.num_beams,
}

for step, batch in enumerate(tqdm(eval_dataloader)):
    with torch.no_grad():
        generated_tokens = accelerator.unwrap_model(model).generate(
            batch["input_ids"],
            attention_mask=batch["attention_mask"],
            **gen_kwargs,
            # top_k=100,
            # top_p=0.9,
            # do_sample=True,
            no_repeat_ngram_size=3,
            length_penalty=1.0,
            num_return_sequences=5,
            # temperature=0.3
        )
```

Result:

```
{
  "rouge-1": {
    "f": 0.2765556796885529,
    "p": 0.2840055329476194,
    "r": 0.2822207147691371
  },
  "rouge-2": {
    "f": 0.11531495092447706,
    "p": 0.11870296027625428,
    "r": 0.11783350736967611
  },
  "rouge-l": {
    "f": 0.2481989948995724,
    "p": 0.2569447884469682,
    "r": 0.2508616852042255
  }
}
```

- Beams\_number (1) == use greedy + ngram(2)

Epoch=20

Input size = 2048

Output size = 128

```
gen_kwargs = {
    "max_length": args.val_max_target_length if args is not None else config.max_length,
    "num_beams": args.num_beams,
}

for step, batch in enumerate(tqdm(eval_dataloader)):
    with torch.no_grad():
        generated_tokens = accelerator.unwrap_model(model).generate(
            batch["input_ids"],
            attention_mask=batch["attention_mask"],
            **gen_kwargs,
            no_repeat_ngram_size=2,
            length_penalty=1.0,
            # num_return_sequences=5,
            # temperature=0.8
        )
```

Result:

```
{
  "rouge-1": {
    "f": 0.25963228071846917,
    "p": 0.26944781520616207,
    "r": 0.2619231585166096
  },
  "rouge-2": {
    "f": 0.09855605665930152,
    "p": 0.10231209321021209,
    "r": 0.09995070765347569
  },
  "rouge-l": {
    "f": 0.2304594506771392,
    "p": 0.24124622876238538,
    "r": 0.23018201278218925
  }
}
```

➤ What is your final generation strategy? (you can combine any of them)

- Beam search (num=7)
- Length penalty =1.0      target 長度限制
- No\_repeat\_ngram\_size=3      target 相同詞彙出現限制
- num\_return\_sequences=7      predict 返回最佳前五句

```
{  
  "rouge-1": {  
    "f": 0.2783196093622653,  
    "p": 0.2854129922686215,  
    "r": 0.2843418305349766  
  },  
  "rouge-2": {  
    "f": 0.1164916516785287,  
    "p": 0.11983329519416126,  
    "r": 0.11913228357826665  
  },  
  "rouge-l": {  
    "f": 0.24909060026246263,  
    "p": 0.2572848723185751,  
    "r": 0.2522693397227175  
  }  
}
```